Influential points, ANOVA for regression and multiple regression

Overview

Review of inference for simple linear regression

Examining unusual data points

Analysis of variance for regression

If there is time: Multiple regression
- Basic ideas
- Nested model comparison
- Related sampling and multiple regression coefficients

Quick review of simple linear regression

The process of building regression models

Choose the form of the model
- Identify and transform explanatory and response variables

Fit the model to the data
- Estimate model parameters

Assess how well the model describes the data
- Analyze the residuals, evaluate unusual points, etc.

Use the model to address questions of interest
- Make predictions, explore relationships, etc.

All models are wrong, but some models are useful

Simple linear regression concepts

Theoretical model: $Y = \beta_0 + \beta_1 x + \epsilon$

Estimated model: $\hat{y} = \hat{\beta}_0 + \hat{\beta}_1 x$

Inference for simple linear regression models
- Hypothesis tests for intercept and slope
- Confidence intervals for slope and line; prediction intervals

Regression diagnostics
- Linearity, Independence, Normality, Equal variance of errors

Hypothesis test for regression coefficients

We can run hypothesis tests to assess whether there is a relationship between y and x, and calculate p-values

- H_0: $\beta_1 = 0$ (slope is 0, so no relationship between x and y
- H_A: $\beta_1 \neq 0$

One type of hypothesis test we can run is based on a t-statistic: $t = \dfrac{\hat{\beta}_1 - 0}{\hat{SE}_{\hat{\beta}_1}}$

- The t-statistic comes from a t-distribution with n - 2 degrees of freedom

$$\hat{SE}_{\hat{\beta}_1} = \frac{\hat{\sigma}_\epsilon}{\sqrt{\sum_{i=1}^{n}(x_i - \bar{x})^2}} \qquad\qquad \hat{SE}_{\hat{\beta}_0} = \hat{\sigma}_\epsilon \sqrt{\frac{1}{n} + \frac{\bar{x}^2}{\sum_{i=1}^{n}(x_i - \bar{x})^2}}$$

Confidence and prediction intervals

1. CI for slope β_1

$$\hat{\beta}_1 \pm t^* \cdot SE_{\hat{\beta}_1} \qquad SE_{\hat{\beta}_1} = \sigma_\epsilon \sqrt{\frac{1}{\Sigma_{i=1}^n (x_i - \bar{x})^2}}$$

$$\beta_1$$

2. CI for regression line μ_Y at point x*

$$\hat{y} \pm t^* \cdot SE_{\hat{\mu}} \qquad SE_{\hat{\mu}} = \sigma_\epsilon \sqrt{\frac{1}{n} + \frac{(x^* - \bar{x})^2}{\Sigma_{i=1}^n (x_i - \bar{x})^2}}$$

3. Prediction interval y

$$\hat{y} \pm t^* \cdot SE_{\hat{y}} \qquad SE_{\hat{y}} = \sigma_\epsilon \sqrt{1 + \frac{1}{n} + \frac{(x^* - \bar{x})^2}{\Sigma_{i=1}^n (x_i - \bar{x})^2}}$$

Regression diagnostics

Linearity, Independence, Normality, Equal variance of errors

Questions?

Statistics for unusual observations

Statistics for unusual observations

There are statistics that are useful for flagging usual observations

- **Outliers (large residuals):** unusual **y** values
- **High leverage points**: usual **x** values
- **Influential points**: both an outlier and a high leverage

Unusual observations can indicate:

- An error in data processing
- A need to modify the model
- An interesting phenomenon

Votes for Buchanan vs population size
2000 presidential election

Palm Beach county

Unusual observations **can also have a big effect on the model fit**

- E.g., a big effect on $\hat{\beta}_0$ $\hat{\beta}_1$

Leverage

High leverage points are predictors **x** that are far from the mean

We can calculate the leverage a data point has using the statistic:

$$h_i = \frac{1}{n} + \frac{(x_i - \bar{x})^2}{\sum_{j=1}^{n}(x_j - \bar{x})^2}$$

High leverage points can have a big impact on the model that is fit!!!

R: hatvalues()

$$\Sigma_{i=1}^{n} h_i = 2$$

Typical: $h_i = 2/n$
High: $h_i = 4/n$
Very high: $h_i = 6/n$

Outliers: standardized residuals

The **standardized residual** for the i[th] data point in a regression model can be computed using:

$$stdres_i \; = \; \frac{y_i - \hat{y}}{\hat{\sigma}_\epsilon \sqrt{1 - h_i}}$$

Puts residuals on a 'normalized' scale

Makes residuals at the ends a bit larger to deal with the fact that they are 'overfit'

R: rstandard()

Relationship between log endowment and salary

Salary ($)

Log endowment size (log($))

Outliers: studentized residuals

The **studentized residual** for the i[th] data point in a regression model can be computed using:

$$studres_i \;=\; \frac{y_i - \hat{y}}{\hat{\sigma}_{(i)}\sqrt{1 - h_i}}$$

Votes for Buchanan vs population size
2000 presidential election

Palm Beach county

Number of votes for Buchanan

County population size

Here $\hat{\sigma}_{(i)}$ is the an estimate of $\hat{\sigma}_\epsilon$

with the i[th] point removed

Q: Why might we want to remove the i[th] point when calculating $\hat{\sigma}_\epsilon$?

A: Outliers could have a big effect on our estimate of $\hat{\sigma}_\epsilon$

R: rstudent ()
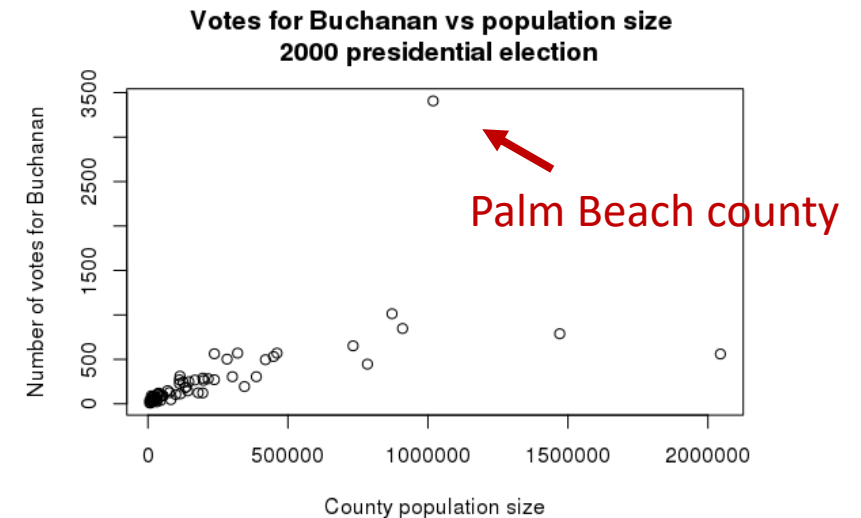
Cook's distance

The amount of influence a point has on a regression line depends on:
- The size of the residual e_i
- The amount of leverage h_i

Cook's distance is a statistic that captures how much influence a point has on a regression line

$$D_i \;=\; \frac{(stdres_i)^2}{k+1}\,\frac{h_i}{1-h_i}$$

Larger for larger
residuals (outliers)

Larger for high
leverage points

R: cooks.distance ()

Where *k* is the number of predictors in the model
- For simple linear regression k = 1   (just a single predictor x)

Cook's distance

The amount of influence a point has on a regression line depends on:
- The size of the residual e_i
- The amount of leverage h_i

Cook's distance is a statistic that captures how much influence a point has on a regression line

$$D_i \;=\; \frac{(stdres_i)^2}{k+1}\,\frac{h_i}{1-h_i}$$

Larger for larger
residuals (outliers)

Larger for high
leverage points

Rule of thumb:
- Moderately influential: $D_i > 0.5$
- Very influential: $D_i > 1$

R: cooks.distance ()

Cook's distances for salary ~ $\log_{10}(\text{endowment})$

plot(lm_fit, 4)

Unusual points rules of thumb

Statistic	Moderately unusual	Very unusual
Leverage, h_i	Above $2(k + 1)/n$	Above $3(k + 1)/n$
Standardized residual	Beyond ± 2	Beyond ± 3
Studentized residual	Beyond ± 2	Beyond ± 3
Cook's D	Above 0.5	Above 1.0

Where:
- k is the number of explanatory variables
- n is the number of data points

From STAT2 by Cannon, Cobb, Hartlaub, Legler, Lock, Moore, Rossman, and Witmer

Questions?

Let's try it in R...

Analysis of Variance (ANOVA) for regression

Analysis of Variance (ANOVA) for regression

In an analysis of variance, we break down the **total variability** in the **response variable y** into:

- 1. the variability explained by the model
- 2. the variability not explained by the model
  - i.e., the residuals

Large

Amount of
statistics
geekiness

Small

Weak   Strong

**Desire to make a graph
representing a linear correlation**

Analysis of Variance (ANOVA) for regression

In an analysis of variance, we break down the **total variability** in the **response variable y** into:

- 1. the variability explained by the model

- 2. the variability not explained by the model

  - i.e., the residuals

We can express this as:

- SSTotal = SSModel + SSResidual

$$y - \bar{y} = (\hat{y} - \bar{y}) + (y - \hat{y})$$

Added and subtracted $\hat{y}$

This equal 0   (proof via algebra)

$$\sum_{i=1}^{n}(y_i - \bar{y})^2 = \sum_{i=1}^{n}(\hat{y}_i - \bar{y})^2 + (y_i - \hat{y}_i)^2 + \cancel{2(y_i - \hat{y}_i)(\hat{y}_i - \bar{y})}$$

SSResidual

SSTotal

SSModel

$\bar{y}$

$\hat{y}$

The coefficient of determination r^2

The **percentage of the total variability explained by the model** is given by

$$r^2 = \frac{SSModel}{SSTotal} = 1 - \frac{SSResidual}{SSTotal}$$

SSResidual

SSTotal

SSModel

$\bar{y}$

$\hat{y}$

We can express this as:

- SSTotal = SSModel + SSResidual

$$y - \bar{y} = (\hat{y} - \bar{y}) + (y - \hat{y})$$

Added and subtracted $\hat{y}$

This equal 0   (proof via algebra)

$$\sum_{i=1}^{n}(y_i - \bar{y})^2 = \sum_{i=1}^{n}(\hat{y}_i - \bar{y})^2 + (y_i - \hat{y}_i)^2 + 2(y_i - \hat{y}_i)(\hat{y}_i - \bar{y})$$

Hypothesis test based on ANOVA for regression

$$F = \frac{SSModel/df_{model}}{SSResidual/df_{error}}$$

$df_{model} = 1$

$df_{error} = n - 2$

If the null hypothesis is true that $\beta_1 = 0$:

- Both the numerator and denominator are estimates of σ^2
- F comes from an F-distribution with df_{model}, df_{error} degrees of freedom

- For simple linear regression, this gives the same results as running a t-test.
  - $F = t^2$

Analysis of Variance (ANOVA) for regression in R

You can create an ANOVA table for regression relationships in R using:

- anova(lm_fit)

SSModel

SSResidual

F

```
lm_fit <- lm(salary_tot ~ log_endowment, data = assistant_data)

anova(lm_fit)
|
```

```
Analysis of Variance Table

Response: salary_tot
                  Df        Sum Sq      Mean Sq    F value             Pr(>F)
log_endowment      1  132879258586  132879258586   764.29  < 0.0000000000000022 ***
Residuals       1173  203936190958     173858645
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Analysis of Variance (ANOVA) for regression in R

You can create an ANOVA table for regression relationships in R using:

- anova(lm_fit)

We can check that the ANOVA relationships holds: SSTotal = SSModel + SSResidual using:

- The original data y values
- lm_fit$residuals
- lm_fit$fitted.values

You can also check that F = t^2 by comparing anova(lm_fit) and summary(lm_fit) values

Homework 7!

Multiple regression

Multiple regression

In multiple regression we try to predict a quantitative response variable y using several predictor variables $x_1, x_2, ... , x_k$

For multiple linear regression, the underlying model is:

$$y = \beta_0 + \beta_1 \cdot x_1 + \beta_2 \cdot x_2 + ... \beta_k \cdot x_k + \epsilon$$

We estimate coefficients using a data set to make predictions $\hat{y}$

$$\hat{y} = \hat{\beta}_0 + \hat{\beta}_1 \cdot x_1 + \hat{\beta}_2 \cdot x_2 + ... + \hat{\beta}_k \cdot x_k$$

Multiple regression

$$\hat{y} \;=\; \hat{\beta}_0 \;+\; \hat{\beta}_1 \cdot x_1 \;+\; \hat{\beta}_2 \cdot x_2 \;+\; ... \;+\; \hat{\beta}_k \cdot x_k$$

There are many uses for multiple regression models including:

- To make predictions as accurately as possible

- To understand which predictors (x) are related to the response variable (y)

Multiple regression

$$\text{salary} \ = \ \hat{\beta}_0 \ + \ \hat{\beta}_1 \cdot \text{f(endowment)} \ + \ \hat{\beta}_2 \cdot \text{g(enrollment)}$$

Let's explore this in R...

Nested model comparison

We can also assess whether a particular subset of *q* parameters is 0

$$H_0: \quad \beta_h = \beta_i = \ldots = \beta_g = 0$$

To do this we:

1. Fit the model without these features
2. Calculate the $SSRes_{Reduced}$ for the model without these predictors
3. Compare it to the full model $SSRes_{Full}$ with an F-statistic:

$$F = \frac{(SSRes_{Reduced} - SSRes_{Full})/q}{SSRes_{Full}/(n-k-1)}$$

where q is the number of additional terms in the full model

F PDF (One-Sided Test at Alpha = 0.05)

df_1 = $df_{Reduced}$ - df_{Full}
df_2 = df_{Full}

Relating simple and multiple regression

Suppose we fit both a simple and multiple regression models to the same data.

simple linear regression coefficient

Simple regression model: $\hat{y} = \hat{\beta}_{0(1)} + \hat{\beta}_{1(1)} \cdot x_1$

multiple linear regression coefficient

Multiple regression model: $\hat{y} = \hat{\beta}_{0(2)} + \hat{\beta}_{1(2)} \cdot x_1 + \hat{\beta}_{2(2)} \cdot x_2$

Question: How are the coefficients $\hat{\beta}_{1(1)}$ and $\hat{\beta}_{1(2)}$ related?

Relating simple and multiple regression

Question: How are the simple regression coefficients $\hat{\beta}_{1(1)}$ and the multiple regression coefficient $\hat{\beta}_{1(2)}$ (for a predictor x_1) related?

We can view the multiple regression coefficient $\hat{\beta}_{1(2)}$ as the change in y when we **set the predictions x_2, to a fixed value**.

We can view the simple regression coefficient $\hat{\beta}_{1(1)}$ as the change in y **when we let the other predictor x_2 change with the value of x_1.**

Relating simple and multiple regression

If the predictor x_1 is correlated with x_2, then changing x_1 will be associated with changes in x_2 which in turn will be associated with changes in y.

We can assess the association between x_1 and x_2, using regression:

$$x_2 \ = \ \hat{\delta}_0 \ + \ \hat{\delta}_1 \cdot x_1$$

We can then relate the change in *y* with the change in x_1 in the simple regression coefficient to the multiple regression coefficients as:

$$\hat{\beta}_{1(1)} \cdot x_1 \ = \ \hat{\beta}_{1(2)} \cdot x_1 \ + \ \hat{\beta}_{2(2)} \cdot \hat{\delta}_1 \cdot x_1$$

Let's explore this in R...